

Support Documentation

Can Wang Gordon Leung Jiashi Chen
Jincheng Song Jirathip Kunkanjanathorn

Contents

Introduction	1
Documentation Page	1
Prerequisites	1
Installation	1
Local Development	1
Deployment	1
Data Workflows	2
Prerequisites	2
Data Collection	2
Data Processing	2
Data Ingestion	2
Backend	3
How to run the system	3
How to shutdown the system	3
How to test the system	3
Unit test using pytest with Docker locally	3
Unit test the system with Github Action	3
Test the system with Postman	3
How to deploy the system	4
Build and Push Images to Dockerhub with Github Action	4
Setting up GCP Compute Instance	4
Deploying on GCP	4
Actions for Developers:	8
How to add table to the database	8
How to backup and restore the database	9
Setup Automated Backups	9
Restore from fresh database	9
How to launch new service in the backend	10
How to monitor the system	11
What to do when the system is down	11
Frontend	12
How to run the system	12
How to shutdown the system	12
How to deploy the system	12
Deployment	12
Connect frontend to the backend and APIs	13
Training and knowledge needed to operate the system	13
Database Management and SQL	13
Python Development	13

React Native Development	14
General Development	14
API and Web Frameworks	15
External Services and Data integrations	15
Natural Language Processing and AI	15
Containerization and Deployment	15
General	15
Resources	15
Internal	16
External	16

Introduction

Documentation Page

Prerequisites

- Node.js and npm (Node Package Manager) should be installed on the local machine.
- A basic understanding of the project's structure and documentation requirements
- Access to the project repository and clone the Git repository using the command:

```
git clone https://github.com/topmello/topmello.github.io.git
```

- Or Downloading and extracting the ZIP file of the repository.

Installation

Upon cloning or when updating dependencies, execute:

```
npm install
```

This command will ensure all necessary dependencies are installed.

Local Development

Whether maintaining the documentation or familiarizing oneself as a new member, it's crucial to preview changes locally. Start the local development server with:

```
npm run start
```

This server will launch a browser window where most modifications will be reflected live without needing to restart the server.

Deployment

Ensure membership in the organization. If not a member, fork the repository and proceed to deploy on the forked version.

Replace <Your GitHub username> with the respective GitHub username:

```
GIT_USER=<Your GitHub username> npm run deploy
```

Data Workflows

Prerequisites

- Python 3.8 or above and Pip should be installed on the local machine.
- Access to to the project repository and clone the Git repository using the command: “`bash git clone https://github.com/topmello/settle-aid-data-wrangling.git`”
- Or downloading and extracting the ZIP file of the repository.

Data Collection

In order to collect the open data, the following datasets are accessed from the Melbourne Open Data Platform using CSV format which require to download manually:

Name	Frequency of Updates	License
Landmarks and places of interest	Monthly	CC BY
Business establishments location and industry classification	Annually	CC BY
Café, restaurant, bistro seats	Annually	CC BY

Data Processing

After the data is collected, the data should be store in the `./data` folder in the project directory. Then, the data processing can be done by running each scipt related to the location type including restaurant, landmark, grocery, and pharmacy.

Noted that this step will take a long time to finish and it can be skipped to run backend server directly if the data is not updated and already processed.

Data Ingestion

The processed data is stored in the `./data` folder in the project directory as JSON file. The data is then ingested into the database by copying to JSON file to backend repository and running the ingestion script.

Backend

How to run the system

1. Clone the repository or from the zip file
2. Install Docker with Docker Compose
3. Run `docker-compose up -d` in the root directory (Compose V2 do not have - between docker-compose)
4. SSH into the backend container with `docker exec -it backend bash`
5. Migrate the database with `alembic upgrade head` inside the container (For the first time or when there's a change in the database schema)
6. Insert the data with `python -m scripts.insert_data` inside the container (For the first time or when there's a change in the dataset)

How to shutdown the system

1. Run `docker-compose down` in the root directory

How to test the system

Unit test using pytest with Docker locally

1. Run `docker-compose -d` in the root directory
2. Run `docker exec -it backend pytest` to run the test

Noted that this test will remove all the data in the database where the test is run. The better way to test is to use CI/CD with Github Action.

Unit test the system with Github Action

This workflow is triggered whenever there's a push to the main branch.

1. Checks out repository.
2. Builds and starts Docker Compose services using the dev configuration.
3. Runs pytest within backend service.
4. Shuts down and removes the containers afterward.

Test the system with Postman

1. Ensure the system is up and running by following the provided steps under "How to run the system".
2. Launch the Postman application on your machine.
3. In the Postman interface, select the desired HTTP method (e.g., GET, POST, PUT, DELETE) from the dropdown.
4. Enter the system's endpoint URL you wish to test in the request URL field.
5. Depending on the endpoint:
 - Add necessary headers by clicking on the "Headers" tab and entering key-value pairs.

- For methods like POST or PUT, click on the “Body” tab and input the required data in the appropriate format (e.g., JSON, form data).
 - If necessary, add query parameters by clicking on the “Params” tab and inputting key-value pairs.
6. Click the “Send” button to initiate the request to the backend.
 7. Inspect the response received in the Postman response section below. Here, you can view status codes, response times, returned data, or any error messages.

How to deploy the system

Build and Push Images to Dockerhub with Github Action

Activated either manually or when there’s a push to the deploy branch. The steps include:

1. Checking out repository.
2. Setting up QEMU & Docker Buildx.
3. Logging into Docker Hub using saved credentials.
4. Building Docker images from Dockerfiles (Dockerfile.db & Dockerfile.backend).
5. Pushing these images to Docker Hub.

Setting up GCP Compute Instance

Before deploying on GCP, ensure the VM instance ready. To set up a VM instance:

1. Go to the GCP Console at <https://console.cloud.google.com/>.
2. Navigate to the Compute Engine and then VM Instances.
3. Click on “Create Instance.”
4. Fill out the necessary details like Name, Region, Zone, Machine type, etc.
5. In the Boot Disk section, select an Ubuntu as OS.
6. Under the Firewall settings, make sure to allow HTTP and HTTPS traffic if your application needs to be accessed over the internet.
7. Once filled out, click “Create” to instantiate your VM.
8. SSH into the instance and install Docker and Docker Compose. The instructions can be found here: <https://docs.docker.com/engine/install/ubuntu/>
9. Install VIM using `sudo apt install vim`. (Optional)

Deploying on GCP

1. SSH into GCP Instance: `gcloud compute ssh <instance-name> --zone <zone>`
2. Change directory: `cd ..` (Optional)
3. Make sure `docker-compose.yaml` and `<domain-name>` is exist in the directory. And, the domain DNS has been point to the backend virtual machine

external IP address. For the first time configuration files for production are needed to be created as below. Noted that the keys and domain name needed to be changed according to the new production environment.

- docker-compose.yaml

```
version: '3'
services:
  db:
    container_name: settle-aid-db
    image: jirathipk/postgres-vec-geo:latest
    restart: always
    environment:
      - POSTGRES_DB=database
      - POSTGRES_USER=db_user
      - POSTGRES_PASSWORD=password1234
    volumes:
      - database_volume:/var/lib/postgresql/data/
      - dbbackups_volume:/backups
  redis:
    container_name: settle-aid-redis
    image: redis:latest
    restart: always
    command: redis-server --requirepass topmelloredis --loglevel verbose
    volumes:
      - redis_volume:/data
  backend:
    image: jirathipk/settle-aid-backend:latest
    container_name: settle-aid-backend
    user: myuser
    environment:
      - DATABASE_HOSTNAME=db
      - DATABASE_NAME=database
      - DATABASE_PORT=5432
      - DATABASE_PASSWORD=password1234
      - DATABASE_USERNAME=db_user
      - SECRET_KEY=SECRET_KEY
      # Generate a new key with openssl rand -hex 32
      - REFRESH_SECRET_KEY=REFRESH_SECRET_KEY
      # Generate a new key with openssl rand -hex 32
      - REFRESH_TOKEN_EXPIRE_DAYS=7
      - ALGORITHM=HS256
      - ACCESS_TOKEN_EXPIRE_MINUTES=30
      - MAPBOX_ACCESS_TOKEN={{MAPBOX_ACCESS_TOKEN}}
      # Create an account and get the token from https://account.mapbox.com/
```



```

- DOC_USERNAME=topmello
- DOC_PASSWORD=da7da0df508738e37f18
- REDIS_HOSTNAME=redis
- REDIS_PORT=6379
- REDIS_PASSWORD=topmelloredis
- USER_CACHE_EXPIRY=3600
- TRANSFORMERS_CACHE=/usr/src/app/transformers_cache
- PYTEST_ADDOPTS="-o cache_dir=/usr/src/app/.pytest_cache"
- VIRTUAL_HOST=staging.settle-aid.tech
- VIRTUAL_PORT=8000
- LETSENCRYPT_HOST=staging.settle-aid.tech
- LETSENCRYPT_EMAIL=jirathip.ku@gmail.com
depends_on:
- db
- redis

pgbackups:
  container_name: settle-aid-db-backup
  image: prodrigestivill/postgres-backup-local
  restart: always
  user: postgres:postgres
  volumes:
    - dbbackups_volume:/backups
  links:
    - db
  depends_on:
    - db
  environment:
    - POSTGRES_HOST=db
    - POSTGRES_DB=database
    - POSTGRES_USER=db_user

nginx-proxy:
  image: nginxproxy/nginx-proxy
  container_name: nginx-proxy
  ports:
    - "80:80"
    - "443:443"
  volumes:
    - /var/run/docker.sock:/tmp/docker.sock:ro
    - certs:/etc/nginx/certs
    - vhost:/etc/nginx/vhost.d
    - html:/usr/share/nginx/html
    - ./api.settle-aid.tech:/etc/nginx/vhost.d/api.settle-aid.tech # production domain

acme-companion:

```

```

image: nginxproxy/acme-companion
container_name: nginx-proxy-acme
volumes:
  - /var/run/docker.sock:/var/run/docker.sock:ro
  - certs:/etc/nginx/certs
  - html:/usr/share/nginx/html
  - vhost:/etc/nginx/vhost.d
  - acme:/etc/acme.sh
environment:
  - NGINX_PROXY_CONTAINER=nginx-proxy
depends_on:
  - nginx-proxy

volumes:
  certs:
  vhost:
  html:
  acme:
  database_volume:
  redis_volume:
  models_volume:
  dbbackups_volume:

  • domain name for NGINX proxy

location /logs/stream/ {
    proxy_pass http://localhost:8000;
    proxy_set_header Host $host;
    proxy_set_header X-Real-IP $remote_addr;
    proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;

    # SSE specific configurations
    proxy_http_version 1.1;
    proxy_set_header Connection '';
    proxy_buffering off;
    proxy_cache off;
    send_timeout 600s;
}

location /track-sio/sio/ {
    proxy_pass http://localhost:8000;
    proxy_set_header Host $host;
    proxy_set_header X-Real-IP $remote_addr;
    proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;

    # WebSocket specific configurations

```

```

proxy_http_version 1.1;
proxy_set_header Upgrade $http_upgrade;
proxy_set_header Connection "upgrade";

# Disable buffering when the nginx proxy gets very busy (protects upstream service)
proxy_buffering off;
}

```

4. Pull the Latest Docker Compose Configuration: `sudo docker-compose pull`
5. Start the Containers: `sudo docker-compose -p settle-aid up -d`
 - The `-p` flag is to set a project name, which can be useful for running multiple environments on the same host
 - The `-d` flag is to run the containers in the background
6. For the first time, run the migration script:


```
sudo docker exec -it backend alembic upgrade head
```

```
sudo docker exec -it backend python -m scripts.insert_data
```

Actions for Developers:

- Modifications: If there are modifications or additions to packages, update `requirements.txt` so the Docker build process incorporates these changes.
- GitHub Workflows: The Python application test runs on pushes to the main branch, and the Dockerhub build and push are triggered either manually or when pushing to the deploy branch.

How to add table to the database

1. While inside the backend container, navigate to the directory where your models are defined.
2. Create or modify an ORM model to define the structure of your new table.
3. Generate a new migration script using the command:


```
alembic revision -m "Add new_table_name table"
```
4. Edit the generated migration script in the `versions` directory, ensuring the `upgrade()` method contains logic to create your new table and the `downgrade()` method contains logic to remove it.
5. Apply the migration using


```
alembic upgrade head
```
6. Define ORM models in the `app/models.py` to interact with the new table accordingly.

Note: Please avoid running alembic revision on the production database. Instead, create the migration locally and apply it to the production database.

How to backup and restore the database

Setup Automated Backups

We use the docker-postgres-backup-local image to facilitate our backup tasks. Once the service is started, this will automatically create backups of database daily and maintain.

<https://github.com/prodrigestivill/docker-postgres-backup-local>

Restore from fresh database

To restore a backup to a fresh database:

```
docker exec -it db psql \
--username=db_user \
--dbname=postgres -c "DROP DATABASE database;"

docker exec -it db psql \
--username=db_user \
--dbname=postgres -c "CREATE DATABASE database;"

docker exec -it db psql \
--username=db_user \
--dbname=database -c "CREATE EXTENSION IF NOT EXISTS postgis;"

docker exec -it db psql \
--username=db_user \
--dbname=database -c "CREATE EXTENSION IF NOT EXISTS vector;"

sudo docker exec \
-it db /bin/sh \
-c "zcat /backups/last/database-latest.sql.gz | \
psql --username=db_user --dbname=database -W"
```

Change permission for backups If necessary, adjust the permissions for the backup files.

```
docker exec -u root -it db-backup chown -R 999:999 /backups
```

Check backup volumes content To inspect the contents of the backup volume:

```
sudo docker run \
--rm -it \
-v settle-aid_dbbackups_volume:/volume_content alpine:latest /bin/sh
```

How to launch new service in the backend

1. Pull the Desired Image: If you haven't already, ensure that the desired service's Docker image is available on your system. If it's on a public registry like Docker Hub, you can pull it using:

```
docker pull <image-name>:<tag>
```

Replace `<image-name>:<tag>` with the name and the desired tag/version of the image.

2. Modify the Docker Compose File: Navigate to the directory containing your `docker-compose.yml` file and open it for editing.
3. Add the New Service: In the `docker-compose.yml` file, add a new service definition for the image you've just pulled. For instance:

```
services:
  ...
  new-service-name:
    image: <image-name>:<tag>
    ports:
      - "<external-port>:<internal-port>"
    environment:
      - "ENV_VAR_NAME=value"
    volumes:
      - "/path/on/host:/path/in/container"
    depends_on:
      - "another-service-name"
```

Replace placeholders as appropriate:

- `<image-name>:<tag>` with the image's name and tag.
 - `<external-port>:<internal-port>` to map ports from the container to the host.
 - Adjust `environment` to set any environment variables the service needs.
 - The `volumes` section can be used to mount directories from the host into the container.
 - `depends_on` ensures that the new service starts only after another specified service has started.
4. Launch the New Service: With the service added to the Docker Compose file, navigate to the directory containing the file and run:

```
docker-compose up -d new-service-name
```

This command will start only the new service and any services it depends on. If you want to start all services defined in the Compose file, simply use:

```
docker-compose up -d
```

5. Monitor the Service: You can check the logs of the newly launched service with:

```
docker-compose logs -f new-service-name
```

By following these steps, you'll successfully launch a new service in the backend using Docker Compose. Remember to consult the documentation or README of the specific service image for any particular configurations or environment settings.

How to monitor the system

1. Logging Pages: The developed logging pages are essential tools for observing the system's behavior. Regularly review the `/logs/` endpoint, as it maintains a record of all requests made to the backend. By analysing these logs, system performance, identify patterns, detect anomalies, and troubleshoot issues can be gauged when they arise.
2. Docker Logs: Docker provides built-in logging mechanisms for its containers. You can access the logs of a specific container using the following command:

```
docker-compose logs -f service-name
```
3. System Metrics: Utilise `docker stats` to monitor the resource usage of your containers. This command provides a live stream of container performance metrics such as CPU usage, memory consumption, network IO, and disk IO. By monitoring these metrics, you can gain insights into the resource demands of each service and make informed decisions on system scaling or optimization.

What to do when the system is down

1. Verify the Outage: Ensure it's not just a local or isolated issue.
2. Check Logs: Review application logs for errors or warnings.
3. Check External Services: Make sure dependencies, like databases or third-party APIs, are operational.
4. Restart Services: A simple service restart might solve temporary glitches.

```
sudo docker-compose up -d -p settle-aid
```
5. Review Resources: Ensure the system hasn't run out of essential resources like CPU, RAM, or storage.

Frontend

How to run the system

1. Install Node.js 18 LTS
2. Clone or extract the frontend code to local folder
3. Run commandline tool or open code folder in code editor
4. Run the following command in folder that contains all code files
`npm install`
5. Now you can run the following command to start the development server
`npm run start`
6. After development server finish starting process, you can scan the QR code in the commandline using Expo Go application(Android) or using system camera(iOS)

How to shutdown the system

Close the commandline tool or use “Ctrl + C” shortcut in commandline to terminate the development server.

How to deploy the system

Deployment

- Official Doc: <https://docs.expo.dev/build-reference/apk/>

`expo login`

Login:

If console throw “verify that the path is correct and try again” when using the commands, add “npx” before each command except “npm” one would help.

Install EAS Cli: Run the following command to install EAS Cli.

`npm install --global eas-cli`

Link the code base to project: Make sure the project ID is correct and accessible which can be check in the Expo Dashboard.

`eas init --id 714abc65-7237-4be7-8349-feffeae9f93d`

Build: Run the following command to build the app in expo server.

```
eas build -p android --profile preview
```

Connect frontend to the backend and APIs

In order to connect the frontend to the backend server, environment variables need to be set up. The following steps are for setting up the environment variables in the frontend.

1. Check for URL or IP address of the backend server and the port number. The default port number is 8000. This depends on whether it is development or production environment. Typically, the development environment is localhost:8000 and the production environment is the IP address of the server with port 8000. With HTTPS enabled, the production environment URL is the domain name of the server.
2. In the frontend, navigate to eas.json and make sure the `EXPO_PUBLIC_API_URL` is the backend server URL. For example, `https://api.settle-aid.tech`.
3. For other APIs such as Google places and weather APIs, ...

Training and knowledge needed to operate the system

Database Management and SQL

- **PostgreSQL:** Familiarity with relational database management, including CRUD operations, indexing, and optimization in PostgreSQL.
- **Alembic:** Understanding of database migrations and the ability to apply, revert, or create new migrations using Alembic.
- **Redis:** A grasp on key-value store principles, Redis data structures, and caching strategies.

Python Development

- **SQLAlchemy:** Experience in working with Object Relational Mapping (ORM) to integrate Python applications with databases.
- **Pydantic:** Knowledge of data validation and parsing using Pydantic, ensuring incoming data complies with expected formats.
- **PyTest:** Proficiency in writing and running unit tests using PyTest to ensure the functionality and robustness of the codebase.
- **PyTorch:** Basics of machine learning and deep learning concepts, and how to utilize PyTorch for training and deploying models.

React Native Development

- **JavaScript** Before diving into any JavaScript-based framework or library, it's important to have a solid understanding of core JavaScript concepts.
- **React** Functional Components and Hooks: Know how to use `useState`, `useEffect`, `useContext`, and other built-in hooks.
- **JSX** The syntax used to write UI elements in React.
- **React Native** Understanding of the lifecycle, layout, and styling specific to React Native. Knowledge about the difference between React Native and standard React (for web).
- **React Native Paper** Knowledge about Paper's UI components and how to implement them. Familiarity with theming and customization.
- **React Native Dates** Understanding how to integrate the date selector with forms and other UI elements.
- **Expo** Basics of creating, running, and debugging an Expo project. Understanding of how to install and use various Expo SDKs. Familiarity with the Expo workflow, and publishing.
- **Axios** How to make API calls (GET, POST, etc.). Handling responses and errors.
- **Redux** Understanding of the Flux architecture and how Redux improves upon it. Creation of actions, reducers, and the store. Middleware, async actions (like `redux-thunk`).
- **Redux Persist** Integration with Redux to persist the store. Knowledge about storage engines and their configuration.
- **React i18next** Internationalization concepts. Integration with React for dynamic translations.
- **Expo Location and React Native Maps** Fetching geolocation data. Integrating maps, adding markers, and interacting with maps.
- **Expo Router** Setting up routing for a React Native app. Transition between different screens/components.

General Development

- **Version Control Using Git** Effective tracking and management of your codebase.
- **Debugging** Utilizing debugging tools for React Native and python.
- **Responsive Design** Ensuring app looks and functions well across a range of device sizes.
- **Async Programming** Familiarity with Promises, `async/await` for handling asynchronous operations.

API and Web Frameworks

- **FastAPI:** Mastery of creating, deploying, and managing APIs using FastAPI, coupled with an understanding of asynchronous programming.
- **SlowAPI:** Knowledge about rate limiting and its implementation using SlowAPI to manage request traffic and protect the system.
- **Python SocketIO Server:** Familiarity with WebSockets and real-time bi-directional communication using SocketIO.

External Services and Data integrations

- **Open Data Melbourne:** Awareness of the datasets available and how to integrate and utilize city-specific data for enhanced platform functionality.
- **MapBox, Google Translate & Google Location APIs:** Understanding of external API integrations, request-response patterns, and error handling. Experience in geolocation services, language translation tools, and map rendering is essential.

Natural Language Processing and AI

- **Huggingface's Sentence Transformers:** Grasp on semantic search, sentence embeddings, and the utilization of transformers for better user input understanding.

Containerization and Deployment

- **Docker:** Proficiency in containerizing applications using Docker, understanding of Docker Compose for multi-container applications, and familiarity with container orchestration.
- **GCP (Google Cloud Platform):** Knowledge of deploying and managing Docker images on GCP, understanding GCP's infrastructure, and its security best practices.

General

- **System architecture:** A holistic understanding of how each component of the tech stack interacts with each other, the data flow, and error handling mechanisms.
- **Continuous Integration/Continuous Deployment:** Familiarity with CI/CD processes, especially with tools like GitHub Actions, to automate testing and deployment workflows.

Resources

- GitHub Repository:
 - Frontend

- Backend
- Documentation page
- Data Wrangling

Internal

- Project Documentation Page
- Backend API Documentation
- Backend Logging
- Backend UI for testing

External

- React Native - Fundamental Framework to build native app for Android and iOS in React
- React Native Paper - UI component and theme library for UI consistency accross devices
- React Native Dates - Date selector component for React Native Paper
- Expo Router - Routing between screens
- Axios - API request library
- Redux - Global state management for React app
- Redux Persist - Data persistance for Redux
- React i18next - Internalization for React
- Expo Location - Providing access to Geolocation
- React Native Maps - Cross-platform map component
- Expo Calendar - Provides an API for interacting with the device's system calendars
- Expo KeepAwake - A React component that prevents the screen from sleeping when rendered.
- Expo Localization - A library that provides an interface for native user localization information.
- Expo Print - A library that provides printing functionality for Android and iOS (AirPrint).
- React Native Share - Provide access to system share API
- Socket.io Client - For realtime messaging
- PostgreSQL: Serving as our primary relational database.
- SQLALchemy: A Python-based ORM.
- Alembic: Alembic is a lightweight database migration tool that enables the creation of database schemas and the migration of data.
- Redis: Cache database for storing frequently accessed data and managing TTL for cache expiry.
- FastAPI: Python-based web framework that enables the creation of APIs.
- Pydantic: Pydantic is a Python library that facilitates data validation and parsing. It is used to validate user inputs, ensuring that the data is of the correct type and format.
- SlowAPI: A rate limiting library for Starlette and FastAPI.

- Open Data Melbourne: Leveraging datasets from Melbourne’s open data platform, our platform gains access to a wealth of city-specific information, ranging from infrastructure to cultural landmarks. This data integration enhances the platform’s accuracy and relevance when offering route suggestions or city insights to users.
- Python SocketIO Server: Facilitating real-time communication, especially for location tracking features.
- MapBox, Google Translate & Google Location APIs: These external APIs are integrated to enrich the platform’s functionalities. MapBox assists in generating routes, Google Translate aids in user input translation, and the Google API is pivotal for location search, all culminating in a comprehensive user experience.
- Huggingface’s Sentence Transformers: This AI-powered integration enhances the platform’s semantic search capabilities.
- PyTorch: PyTorch is a Python-based machine learning library that enables the use of deep learning models. It is used to train and deploy the name generator model.
- PyTest: PyTest is a Python testing framework that enables the creation of unit tests. It is used to ensure the quality and integrity of the codebase, especially when new features are added.
- Docker: Docker containers are used to encapsulate the development environment.